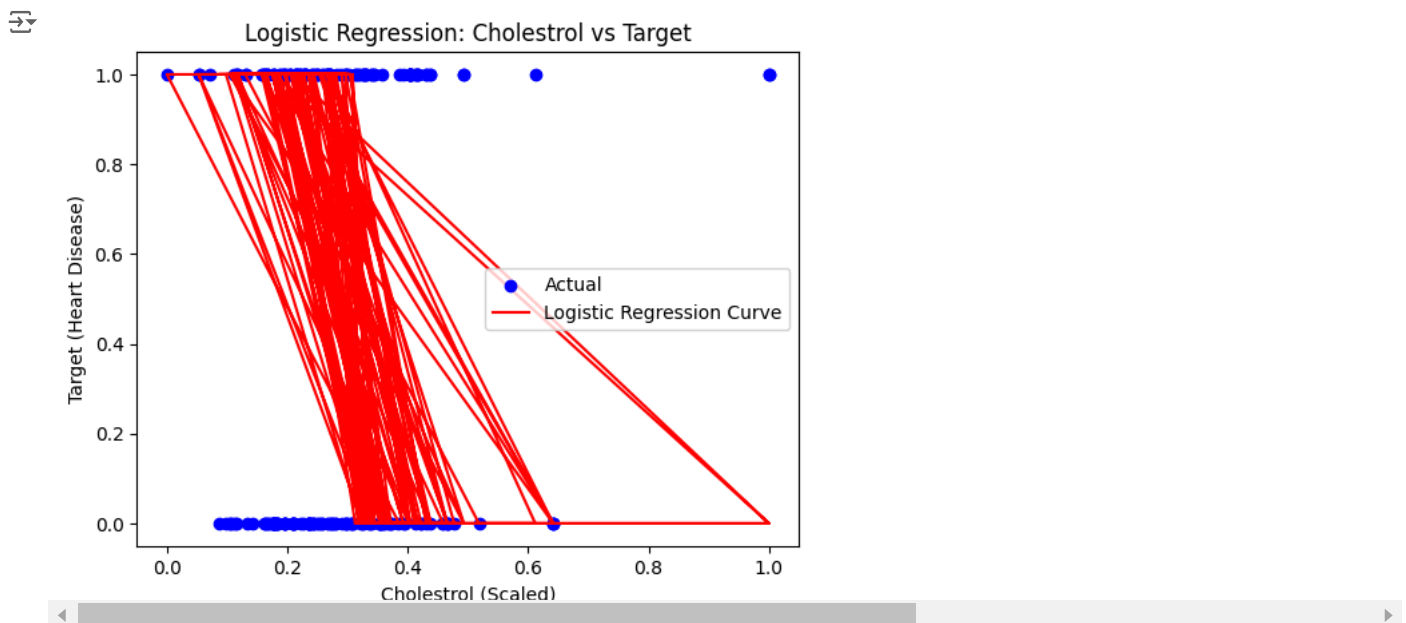


## ✓ Simple Logistic Regression - Heart Disease - Scatter Plot cholesterol vs target

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.preprocessing import MinMaxScaler
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.model_selection import train_test_split
7
8 # Load the dataset
9 data = pd.read_csv('/content/Heart_Disease_Dataset.csv')
10
11 # Extract independent and target variables
12 X = data[['chol']]
13 y = data['target']
14
15 # Min-Max scaling for 'chol'
16 scaler = MinMaxScaler()
17 X_scaled = scaler.fit_transform(X)
18
19 # Split data into training and testing sets
20 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
21
22 # Train a simple logistic regression model
23 model = LogisticRegression()
24 model.fit(X_train, y_train)
25
26 # Get predictions on the test set
27 y_pred = model.predict(X_test)
28
29 # Plot the scatter plot with the regression curve
30 plt.figure()
31 plt.scatter(X_test, y_test, color='blue', label='Actual')
32 plt.plot(X_test, y_pred, color='red', label='Logistic Regression Curve')
33 plt.xlabel('Cholesterol (Scaled)')
34 plt.ylabel('Target (Heart Disease)')
35 plt.title('Logistic Regression: Cholesterol vs Target')
36 plt.legend()
37 plt.show()

```



## ✓ Simple Linear Regression - Car Insurance - Scatter Plot - Insured value vs payment

```


1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.model_selection import train_test_split
5 from sklearn.linear_model import LinearRegression
6 from sklearn.metrics import mean_squared_error, r2_score
7
8 # Load the dataset

```

```

9 data = pd.read_csv('/content/Car_Insurance_Dataset.csv')
10
11 # Extract independent and target variables
12 X = data[['Insured']]
13 y = data['Payment']
14
15 # Split data into training and testing sets
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
17
18 # Train a simple linear regression model
19 model = LinearRegression()
20 model.fit(X_train, y_train)
21
22 # Get predictions on the test set
23 y_pred = model.predict(X_test)
24
25 # Evaluate the model
26 mse = mean_squared_error(y_test, y_pred)
27 r2 = r2_score(y_test, y_pred)
28
29 # Plot the scatter plot with the regression curve
30 plt.figure()
31 plt.scatter(X_test, y_test, color='blue', label='Actual')
32 plt.plot(X_test, y_pred, color='red', label='Linear Regression Curve')
33 plt.xlabel('Insured Value')
34 plt.ylabel('Payment')
35 plt.title('Linear Regression: Insured Value vs Payment')
36 plt.legend()
37 plt.show()
38
39 print(f"Mean Squared Error: {mse}")
40 print(f"R-squared: {r2}")

```



Mean Squared Error: 290225809666.27216

R-squared: 0.827746780422404

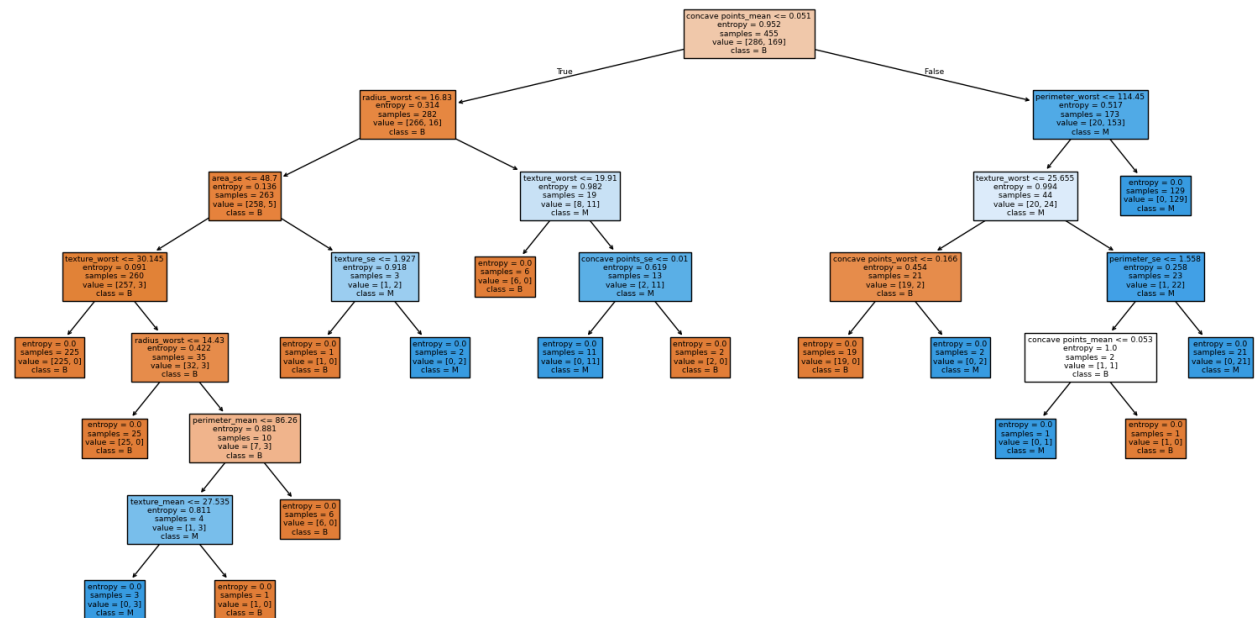
## ✓ Decision Tree - Breast Cancer - Visualize the tree - Histogram

```


1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.model_selection import train_test_split
6 from sklearn.tree import plot_tree
7 from sklearn.preprocessing import LabelEncoder
8
9 # Load the dataset
10 data = pd.read_csv('/content/Breast_Cancer_Dataset.csv')
11
12 # Drop the 'id' column
13 data = data.drop('id', axis=1)
14
15 # Encode the 'diagnosis' column
16 label_encoder = LabelEncoder()
17 data['diagnosis'] = label_encoder.fit_transform(data['diagnosis'])
18 # 'B' will be 0, 'M' will be 1

```

```
19
20 # Separate features (X) and target (y)
21 X = data.drop('diagnosis', axis=1)
22 y = data['diagnosis']
23
24 # Split data into training and testing sets
25 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
26
27 # Create a Decision Tree Classifier with 'entropy' criterion
28 model = DecisionTreeClassifier(criterion='entropy', random_state=42)
29
30 # Train the model
31 model.fit(X_train, y_train)
32
33 # Visualize the decision tree
34 plt.figure(figsize=(20, 10))
35 plot_tree(model, feature_names=X.columns, class_names=['B', 'M'], filled=True)
36 plt.show()
37
38 # Plot the histogram for the classes 'B' and 'M'
39 plt.figure()
40 plt.hist(data[data['diagnosis'] == 0]['diagnosis'], bins=2, alpha=0.5, label='Benign (B)')
41 plt.hist(data[data['diagnosis'] == 1]['diagnosis'], bins=2, alpha=0.5, label='Malignant (M)')
42 plt.xlabel('Diagnosis')
43 plt.ylabel('Frequency')
44 plt.title('Histogram of Diagnosis Classes')
45 plt.legend()
46 plt.show()
```



Histogram of Diagnosis Classes



## ✓ K-means - Iris - Scatter Plot - Petal width vs Length

```

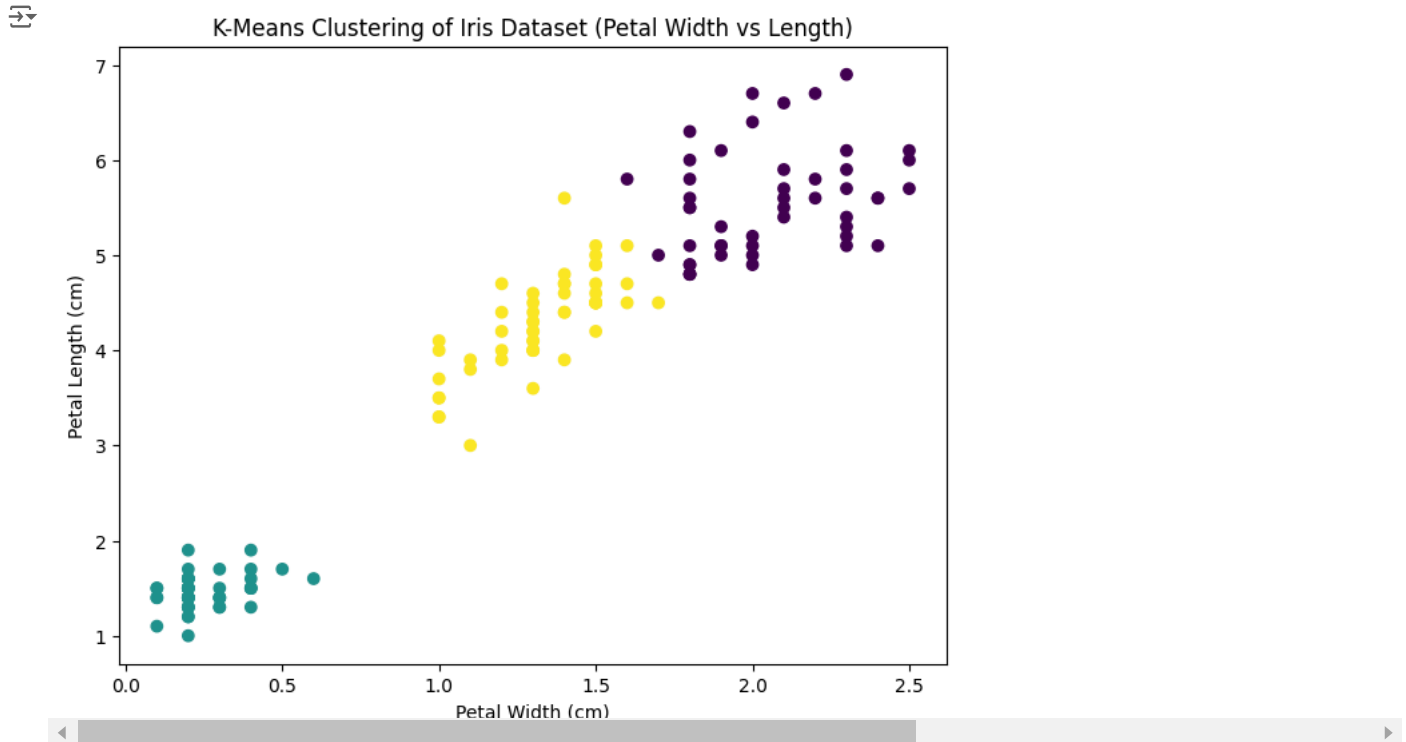
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 from sklearn.cluster import KMeans
4 from sklearn.preprocessing import StandardScaler
5
6 # Load the Iris dataset
7 data = pd.read_csv('/content/Iris_Dataset.csv')
8
9 # Extract features for clustering
10 X = data[['petal.width', 'petal.length']]
11
12 # Standardize the features
13 scaler = StandardScaler()
14 X_scaled = scaler.fit_transform(X)
15
16 # Apply K-Means clustering with 3 clusters
17 kmeans = KMeans(n_clusters=3, random_state=42)
18 kmeans.fit(X_scaled)

```

```

19
20 # Get cluster labels
21 cluster_labels = kmeans.labels_
22
23 # Add cluster labels to the DataFrame
24 data['cluster'] = cluster_labels
25
26 # Plot the scatter plot of Petal Width vs Length, colored by cluster
27 plt.figure(figsize=(8, 6))
28 plt.scatter(data['petal.width'], data['petal.length'], c=data['cluster'], cmap='viridis')
29 plt.xlabel('Petal Width (cm)')
30 plt.ylabel('Petal Length (cm)')
31 plt.title('K-Means Clustering of Iris Dataset (Petal Width vs Length)')
32 plt.show()

```



## ✓ Ensemble - Car Insurance - Random Forest - Box Plot

```


1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.ensemble import RandomForestRegressor
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import mean_squared_error, r2_score
7
8
9 # Load the dataset
10 data = pd.read_csv('/content/Car_Insurance_Dataset.csv')
11
12 # Extract independent and target variables
13 X = data.drop('Payment', axis=1) # Features
14 y = data['Payment'] # Target
15
16 # Split data into training and testing sets
17 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
18
19 # Train a Random Forest Regressor
20 model = RandomForestRegressor(n_estimators=100, random_state=42)
21 model.fit(X_train, y_train)
22
23 # Make predictions on the test set
24 y_pred = model.predict(X_test)
25
26 # Evaluate the model
27 mse = mean_squared_error(y_test, y_pred)
28 r2 = r2_score(y_test, y_pred)
29
30 print(f"Mean Squared Error: {mse}")
31 print(f"R-squared: {r2}")
32
33

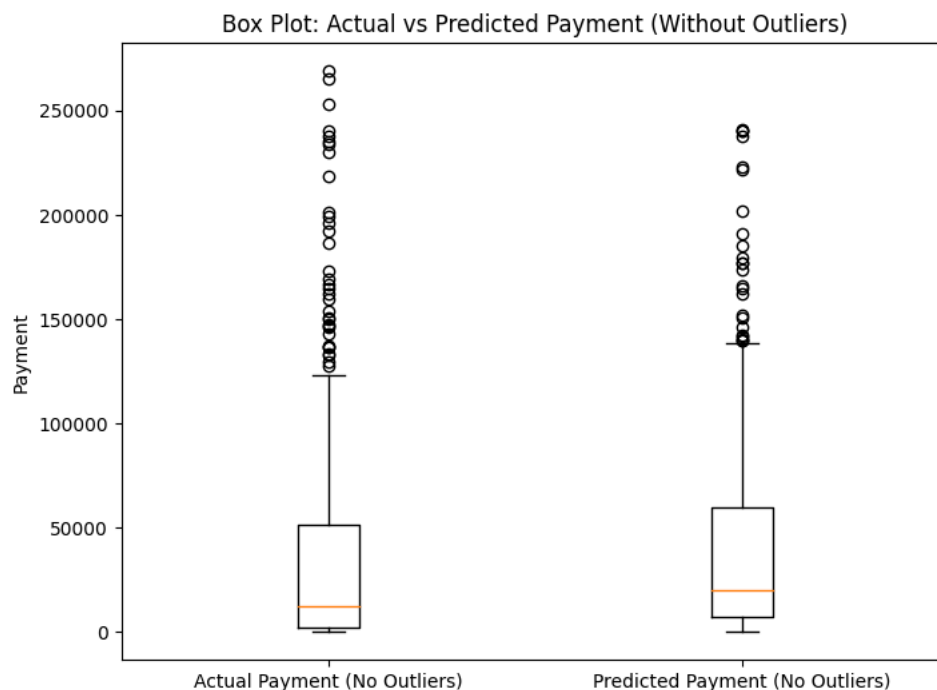
```

```

34 def remove_outliers_iqr(data):
35     Q1 = np.percentile(data, 25)
36     Q3 = np.percentile(data, 75)
37     IQR = Q3 - Q1
38     lower_bound = Q1 - 1.5 * IQR
39     upper_bound = Q3 + 1.5 * IQR
40     data_no_outliers = data[(data >= lower_bound) & (data <= upper_bound)]
41     return data_no_outliers
42
43
44 # Remove outliers from actual and predicted payments
45 y_test_no_outliers = remove_outliers_iqr(y_test)
46 y_pred_no_outliers = remove_outliers_iqr(y_pred)
47
48 # Plot box plots for actual vs predicted values without outliers
49 plt.figure(figsize=(8, 6))
50 plt.boxplot([y_test_no_outliers, y_pred_no_outliers],
51             labels=['Actual Payment (No Outliers)', 'Predicted Payment (No Outliers)'])
52 plt.title('Box Plot: Actual vs Predicted Payment (Without Outliers)')
53 plt.ylabel('Payment')
54 plt.show()

```

 Mean Squared Error: 23865698624.8345  
 R-squared: 0.9866576771943127



## ✓ Multiple Linear Regression - Student Performance - Box Plot

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.linear_model import LinearRegression
5 from sklearn.model_selection import train_test_split
6 from sklearn.metrics import mean_squared_error, r2_score
7
8 # Load the dataset
9 df = pd.read_csv("/content/Student_Performance_Dataset.csv")
10
11 # Define independent and dependent variables
12 X = df[['Previous Scores', 'Hours Studied', 'Sleep Hours']]
13 y = df['Performance Index']
14
15 # Split data into training and testing sets
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
17
18 # Create and train the linear regression model
19 model = LinearRegression()
20 model.fit(X_train, y_train)
21
22 # Make predictions on the test set
23 y_pred = model.predict(X_test)
24
25 # Evaluate the model

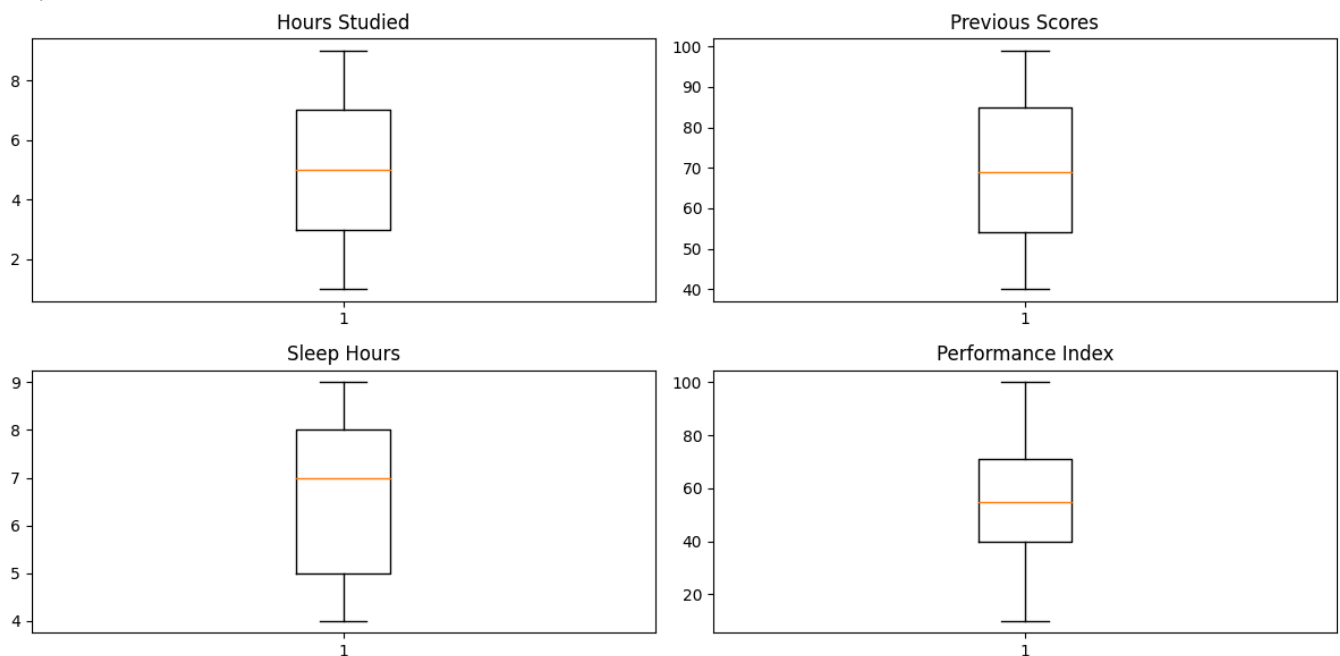
```

```

26 mse = mean_squared_error(y_test, y_pred)
27 r2 = r2_score(y_test, y_pred)
28
29 print(f"Mean Squared Error: {mse}")
30 print(f"R-squared: {r2}")
31
32 # Plot box plots for the specified variables
33 plt.figure(figsize=(12, 6))
34
35 plt.subplot(2, 2, 1)
36 plt.boxplot(df['Hours Studied'])
37 plt.title('Hours Studied')
38
39 plt.subplot(2, 2, 2)
40 plt.boxplot(df['Previous Scores'])
41 plt.title('Previous Scores')
42
43 plt.subplot(2, 2, 3)
44 plt.boxplot(df['Sleep Hours'])
45 plt.title('Sleep Hours')
46
47 plt.subplot(2, 2, 4)
48 plt.boxplot(df['Performance Index'])
49 plt.title('Performance Index')
50
51 plt.tight_layout()
52 plt.show()

```

Mean Squared Error: 4.545107899420576  
R-squared: 0.9877353198949831



## Multiple Logistic Regression - Heart Disease - Scatter plot Resting BP vs Cholesterol

```

1 # Load the dataset
2 df = pd.read_csv("/content/Heart_Disease_Dataset.csv")
3
4 # Define independent and dependent variables
5 X = df[['age', 'trestbps', 'chol']]
6 y = df['target']
7
8 # Split data into training and testing sets
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
10
11 # Create and train the logistic regression model
12 model = LogisticRegression()
13 model.fit(X_train, y_train)
14
15 # Make predictions on the test set
16 y_pred = model.predict(X_test)

```

```


17
18 # Evaluate the model
19 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
20
21 accuracy = accuracy_score(y_test, y_pred)
22 precision = precision_score(y_test, y_pred)
23 recall = recall_score(y_test, y_pred)
24 f1 = f1_score(y_test, y_pred)
25
26 print(f"Accuracy: {accuracy}")
27 print(f"Precision: {precision}")
28 print(f"Recall: {recall}")
29 print(f"F1-score: {f1}")
30
31 # Scatter plot between 'trestbps' and 'chol'
32 plt.figure(figsize=(8, 6))
33 plt.scatter(df['trestbps'], df['chol'], c=df['target'], cmap='viridis')
34 plt.xlabel('Resting Blood Pressure (trestbps)')
35 plt.ylabel('Cholesterol (chol)')
36 plt.title('Scatter Plot: Resting Blood Pressure vs Cholesterol')
37 plt.colorbar(label='Target (Heart Disease)')
38 plt.show()

```

```

↗ Accuracy: 0.6146341463414634
Precision: 0.6224489795918368
Recall: 0.5922330097087378
F1-score: 0.6069651741293532

```



## ✓ Ensemble - XGBoost - Car Insurance - Boxplot

```

1 import pandas as pd
2 import xgboost as xgb
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import mean_squared_error, r2_score
5 import matplotlib.pyplot as plt
6
7 # Load the dataset
8 data = pd.read_csv('/content/Car_Insurance_Dataset.csv')
9
10 # Extract independent and target variables
11 X = data.drop('Payment', axis=1) # Features
12 y = data['Payment'] # Target
13
14 # Split data into training and testing sets
15 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
16
17 # Train an XGBoost Regressor
18 model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=100, random_state=42)
19 model.fit(X_train, y_train)
20
21 # Make predictions on the test set

```



```
22 y_pred = model.predict(X_test)
23
24 # Evaluate the model
25 mse = mean_squared_error(y_test, y_pred)
26 r2 = r2_score(y_test, y_pred)
27
28 print(f"Mean Squared Error: {mse}")
29 print(f"R-squared: {r2}")
30
31 def remove_outliers_iqr(data):
32     Q1 = np.percentile(data, 25)
33     Q3 = np.percentile(data, 75)
34     IQR = Q3 - Q1
35     lower_bound = Q1 - 1.5 * IQR
36     upper_bound = Q3 + 1.5 * IQR
37     data_no_outliers = data[(data >= lower_bound) & (data <= upper_bound)]
38     return data_no_outliers
39
40
41 # Remove outliers from actual and predicted payments
42 y_test_no_outliers = remove_outliers_iqr(y_test)
43 y_pred_no_outliers = remove_outliers_iqr(y_pred)
44
45 # Plot box plots for actual vs predicted values without outliers
46 plt.figure(figsize=(8, 6))
47 plt.boxplot([y_test_no_outliers, y_pred_no_outliers],
48             labels=['Actual Payment (No Outliers)', 'Predicted Payment (No Outliers)'])
49 plt.title('Box Plot: Actual vs Predicted Payment (Without Outliers)')
50 plt.xlabel('Payment')
```